

Recherche visuelle avec Microsoft Foundry et Cohere Embed 4

Recherche visuelle avec Microsoft Foundry et Cohere Embed 4.0



**global AI
community**
Virtuelle Francophone

Recherche visuelle avec Microsoft Foundry et Cohere Embed 4.0

Apprenez à construire des solutions de recherche visuelle en combinant Microsoft Foundry et le modèle Cohere Embed v4. Ce webinaire couvre l'intégration complète du pipeline : **génération d'embeddings vectoriels cohérents pour images et texte, indexation dans Azure AI Search, et requêtes multimodales depuis du texte ou des images.**

Une présentation de ce service à l'aide d'une démonstration sera réalisée lors de ce webinaire.



Serge RETKOWSKY

France AI Global Black Belt
@Microsoft



Paul PETON

AI Architect
@Expertime
AI MVP

LinkedIn Live

jeudi

19/02/2026


à partir de 12h

<https://www.linkedin.com/events/7428033909096996864>

Microsoft Foundry



Microsoft Foundry



Foundry models

demos-resource-project

Search with AI (Ctrl + K)

New Foundry Home Discover Build Operate Docs

Overview

Models

Tools

Solution templates

Models (11307)

View leaderboard Compare models

Search

Sort by Featured

> Collections > Capabilities > Source > Inference tasks > Fine-tuning methods > Industry	claude-sonnet-4-6 Messages gpt-5.2-chat Chat completion, Responses claude-opus-4-6 Messages Kimi-K2.5 Chat completion
	gpt-5.2-codex Responses DeepSeek-V3.2 Chat completion DeepSeek-V3.2-Special Chat completion gpt-5.2 Chat completion, Responses
	Cohere-rerank-v4.0-pro Text classification Cohere-rerank-v4.0-fast Text classification Kimi-K2-Thinking Chat completion gpt-5.1-codex-max Responses
	claude-opus-4-5 Messages gpt-5.1 Chat completion, Responses gpt-5.1-codex Responses
	DeepSeek-V3.1 Chat completion Mistral-Large-3 Chat completion gpt-5-chat Chat completion, Responses claude-haiku-4-5 Messages
	model-router Chat completion claude-opus-4-1 Messages grok-4 Chat completion sora-2 Video generation
	embed-v-4-0 Embeddings, Summarization gpt-5.1-chat Chat completion, Responses grok-4-fast-reasoning Chat completion
	gpt-5-pro Chat completion, Responses Llama-4-Maverick-17B-128E-Instruct-F... Chat completion gpt-5 Chat completion, Responses DeepSeek-V3-0324 Chat completion
	gpt-4.1 Chat completion, Responses gpt-4.1-mini Chat completion, Responses grok-4-fast-non-reasoning Chat completion gpt-4o-transcribe-diarize Speech to text
	FLUX.1-Kontext-pro Text to image, Image to image gpt-5-codex Chat completion, Responses FLUX-1.1-pro Text to image o3 Chat completion, Responses

Modèles Cohere

Microsoft Foundry / demos-resource-project

Overview

Models

Tools

Solution templates

Models (13)

Search

Sort by Featured

Source: Cohere Clear filters

> Collections

> Capabilities

> Source (1)

☐ AI21 Labs 2

☐ Alibaba 1

☐ Anthropic 6

☐ Azure OpenAI 64

☐ Bayer 1

☐ Black Forest Labs 3

☐ BosonAI 1

☐ Bria 1

☒ Cohere 13

embed-v-4-0

Embeddings, Summarization

Cohere-rerank-v4.0-pro

Text classification

Cohere-rerank-v4.0-fast

Text classification

Cohere-rerank-v3.5

Text classification

Cohere-rerank-v3-multilingual

Text classification

Cohere-rerank-v3-english

Text classification

Cohere-command-r

Chat completion

Cohere-command-a

Chat completion

Cohere-command-r-08-2024

Chat completion

Cohere-command-r-plus-08-2024

Chat completion

Cohere-command-r-plus

Chat completion

Cohere-embed-v3-multilingual

Embeddings

Cohere-embed-v3-english

Embeddings

Cohere Embed v 4

Microsoft Foundry

Preview

demos-resource-project

Search with AI (Ctrl + K)

New Foundry

Home

Discover

Build

Operate

Docs

Overview

Models

Tools

Solution templates

embed-v-4-0

Deploy

Details

Deployments

Benchmarks

License

Embed 4 transforms texts and images into numerical vectors

Cohere

Direct from Azure

Version: 6

Cohere's Embed 4 is a multilingual multimodal embedding model. It is capable of transforming different modalities such as images, texts, and interleaved images and texts into a single vector representation. Embed 4 offers state-of-the-art performance across all modalities (texts, images, interleaved texts and image) and in both English and multilingual settings.

Embed 4 supports a 128k context length and an images can have a maximum of 2MM pixels. Embed 4 is capable of vectorizing interleaved texts and images and capturing key visual features from screenshots of PDFs, slides, tables, figures, and more, thereby eliminating the need for complex document parsing. Embed 4 offers a variety of ways for compression both on the number of dimensions and the number-format precision. The model offers byte and binary quantization and matryoshka embeddings for further compression.

Quick facts

Model provider

Cohere

Type

Embeddings, Summarization

Lifecycle

Generally available (GA)

Input type

image, text

Output type

image, text

Context

4096 output

Pricing

View pricing

Model ID

azureml://registries/azureml-cohere/models/embed-v-4-0/versions/6

Others who deployed this model also used

cohere-command-a

Chat completion

Cohere-embed-v3-multilingual

Embeddings


Cohere-embed-v3-english

Embeddings


Cohere Embed v4 — How Embeddings Work

Multimodal embedding model · Text + Images → Vectors · 128K context · Matryoshka dimensions

① Inputs



② Model



③ Output: Embedding Vector



④ Matryoshka Dimensions

(choose precision vs. speed)



⑤ Embedding Types (Compression)



⑥ What You Can Do




Cosine similarity

Cosine Similarity Explained


Measures the angle between two vectors — not their magnitude



$\cos(\theta) \approx 1$
Very Similar



$\cos(\theta) = 0$
No Relationship



$\cos(\theta) = -1$
Completely Opposite

$$\cos(\theta) = (A \cdot B) / (\|A\| \times \|B\|)$$

dot product divided by the product of magnitudes

Microsoft Foundry Cohere Embed 4 demos

1. Recherche visuelle

https://aka.ms/demo_cohere_embed_1

2. Détection des anomalies

https://aka.ms/demo_cohere_embed_2

3. Auto-tagging des images

https://aka.ms/demo_cohere_embed_3

Documentation

- Microsoft :

<https://techcommunity.microsoft.com/blog/azure-ai-foundry-blog/unlock-multi-modal-embed-4-and-multilingual-agentic-rag-with-command-a-on-azure/4404458>

- Model :

<https://ai.azure.com/catalog/models/embed-v-4-0>

<https://ai.azure.com/explore/models/embed-v-4-0/version/1/registry/azureml-cohere?tid=16b3c013-d300-468d-ac64-7eda0820b6d3>



Article Medium

- Article

Building a Multimodal Visual Search Engine with Cohere Embed 4 and Microsoft Foundry



Serge Retkowsky · 7 min read · Dec 11, 2025



3



A complete guide to creating a multilingual fashion search application with text and image queries



 Azure AI Foundry

Prochains évènements



global AI community
Virtuelle Francophone

**Présentation du service
Azure Content Understanding**

Azure Content Understanding est un service Microsoft Foundry qui automatise l'extraction et la structuration de contenu multimodal (documents, images, vidéos, fichiers audios) sans labélisation complexe.

Une présentation de ce service à l'aide de démonstrations sera réalisée lors de ce webinar.



Serge RETKOWSKY
France AI Global Black Belt
@Microsoft



Randa MEZRAG
Consultante DATA & IA
@Finelog-Biseum

LinkedIn Live
jeudi
12/03/2026
à partir de 13h

<https://lnkd.in/ewUig486>



global AI community
Virtuelle Francophone

**Présentation des accélérateurs
Citadel et Spec2Cloud**

Citadel et Spec2Cloud sont deux accélérateurs Microsoft IA. Citadel fournit une **gouvernance centralisée, la conformité et l'observabilité** pour déployer des agents IA sécurisés et contrôlés à l'échelle. Spec2Cloud accélère le **développement d'applications et d'agents IA** via des modèles de spécifications.

Ensemble, ils résolvent le dilemme vitesse vs contrôle : Spec2Cloud vous fait développer vite, Citadel vous assure de le faire en sécurité et conformité. Ces solutions offrent une base pour **transformer rapidement les idées en agents IA gouvernés et production-ready**.



Serge RETKOWSKY
France AI Global Black Belt
@Microsoft



Alibek JAKUPOV
France AI Global Black Belt
@Microsoft



Nicolas ROBERT
Profession
@Cellenza
AI MVP

LinkedIn Live
jeudi
19/03/2026
à partir de 13h

https://lnkd.in/eUe_bBrz

<https://www.linkedin.com/company/global-ai-community-france/posts/?feedView=all>

Replays Global AI Community



[\(7\) Global AI Community Virtuelle Francophone: Posts | LinkedIn](#)

Thank you

